



МИНОБРНАУКИ РОССИИ  
Федеральное государственное бюджетное образовательное  
учреждение  
высшего образования  
**«МИРЭА – Российский технологический университет»**  
**РТУ МИРЭА**

---

ИКБ направление «Киберразведка и противодействие угрозам с  
применением технологий Искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной  
безопасности»

**Отчет по лабораторной работе номер 1**

по дисциплине

«Анализ защищенности систем искусственного интеллекта»

Группа:  
ББМО–02–22  
Коноплич И. С.  
Проверил:  
Спирин А. А.

Москва 2023

Каждый отрывок кода начинается с краткого комментария, объясняющего цель написания этого отрывка.

Результат для каждого отрывка кода для наглядности будет выведен на сером фоне.

```
# Сначала скопируем проект по ссылке в локальную среду Google Colab
!git clone https://github.com/ewatson2/EEL6812_DeepFool_Project.git
Cloning into 'EEL6812_DeepFool_Project'... remote: Enumerating objects: 96, done. remote: Counting
objects: 100% (3/3), done.
remote: Compressing objects: 100% (2/2), done.
remote: Total 96 (delta 2), reused 1 (delta 1), pack-reused 93 Receiving objects: 100% (96/96),
33.99 MiB | 11.54 MiB/s, done. Resolving deltas: 100% (27/27), done.

# Меняем директорию исполнения на папку проекта
%cd /content/EEL6812_DeepFool_Project
/content/EEL6812_DeepFool_Project
# Импортируем необходимые библиотеки
import numpy as np import json, torch
from torch.utils.data import DataLoader, random_split from torchvision import datasets,
models
from torchvision.transforms import transforms

# Импортируем вспомогательные библиотеки из файлов проекта
from models.project_models import FC_500_150, LeNet_CIFAR, LeNet_MNIST,
Net
from utils.project_utils import get_clip_bounds, evaluate_attack, display_attack

# Установим случайное случайное значение в виде переменной rand_seed={"Порядковый номер
ученика группы в Гугл-таблице"}, укажем значение для np.random.seed и torch.manual_seed
rand_seed = 22 np.random.seed(rand_seed) torch.manual_seed(rand_seed)
<torch._C.Generator at 0x7ce54824a550>
# Используем в качестве устройства видеокарту
use_cuda = torch.cuda.is_available()
device = torch.device('cuda' if use_cuda else 'cpu')
# Загрузим датасет MNIST с параметрами mnist_mean = 0.5, mnist_std = 0.5, mnist_dim = 28
mnist_mean = 0.5
mnist_std = 0.5
mnist_dim = 28

mnist_min, mnist_max = get_clip_bounds(mnist_mean,
                                       mnist_std,
                                       mnist_dim)
mnist_min = mnist_min.to(device) mnist_max = mnist_max.to(device)

mnist_tf = transforms.Compose([ transforms.ToTensor(), transforms.Normalize(
    mean=mnist_mean, std=mnist_std)])
mnist_tf_train = transforms.Compose([ transforms.RandomHorizontalFlip(),
    transforms.ToTensor(), transforms.Normalize(
    mean=mnist_mean, std=mnist_std)])
mnist_tf_inv = transforms.Compose([ transforms.Normalize(
    mean=0.0,
    std=np.divide(1.0, mnist_std)), transforms.Normalize(
    mean=np.multiply(-1.0, mnist_std), std=1.0)])
mnist_temp = datasets.MNIST(root='datasets/mnist', train=True,
                           download=True, transform=mnist_tf_train) mnist_train, mnist_val
= random_split(mnist_temp, [50000, 10000])

mnist_test = datasets.MNIST(root='datasets/mnist', train=False,
                           download=True, transform=mnist_tf)

cifar_classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog',
                 'frog', 'horse', 'ship', 'truck']
```

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz Downloading
http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to datasets/mnist/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 122629555.19it/s]
Extracting datasets/mnist/MNIST/raw/train-images-idx3-ubyte.gz to datasets/mnist/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz Downloading
http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to datasets/mnist/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 39677593.78it/s]
Extracting datasets/mnist/MNIST/raw/train-labels-idx1-ubyte.gz to datasets/mnist/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz Downloading
http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to datasets/mnist/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 47687909.56it/s]
Extracting datasets/mnist/MNIST/raw/t10k-images-idx3-ubyte.gz to datasets/mnist/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz Downloading
http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to datasets/mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 19012503.76it/s]
Extracting datasets/mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz to datasets/mnist/MNIST/raw

# Загрузим датасет CIFAR-10 с параметрами cifar_mean = [0.491, 0.482, 0.447] cifar_std =
[0.202, 0.199, 0.201] cifar_dim = 32
cifar_mean = [0.491, 0.482, 0.447]
cifar_std = [0.202, 0.199, 0.201]
cifar_dim = 32
cifar_min, cifar_max = get_clip_bounds(cifar_mean,
                                     cifar_std, cifar_dim)

cifar_min = cifar_min.to(device) cifar_max = cifar_max.to(device)

cifar_tf = transforms.Compose([ transforms.ToTensor(), transforms.Normalize(
    mean=cifar_mean, std=cifar_std)])
cifar_tf_train = transforms.Compose([ transforms.RandomCrop(
    size=cifar_dim, padding=4),
    transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.Normalize(
    mean=cifar_mean, std=cifar_std)])

cifar_tf_inv = transforms.Compose([ transforms.Normalize(
    mean=[0.0, 0.0, 0.0],
    std=np.divide(1.0, cifar_std)), transforms.Normalize(
    mean=np.multiply(-1.0, cifar_mean), std=[1.0, 1.0, 1.0])])

cifar_temp = datasets.CIFAR10(root='datasets/cifar-10', train=True,
                              download=True, transform=cifar_tf_train) cifar_train,
cifar_val = random_split(cifar_temp, [40000, 10000])

cifar_test = datasets.CIFAR10(root='datasets/cifar-10', train=False,
                              download=True, transform=cifar_tf)

Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to datasets/cifar-10/cifar-10-python.tar.gz
100%|██████████| 170498071/170498071 [00:03<00:00, 43315076.71it/s]
Extracting datasets/cifar-10/cifar-10-python.tar.gz to datasets/cifar-10 Files already downloaded
and verified

# Выполним настройку и загрузку DataLoader batch_size = 64 workers = 4
batch_size = 64
workers = 4
mnist_loader_train = DataLoader(mnist_train, batch_size=batch_size,
                               shuffle=True, num_workers=workers) mnist_loader_val =
DataLoader(mnist_val, batch_size=batch_size,
           shuffle=False, num_workers=workers) mnist_loader_test =
DataLoader(mnist_test, batch_size=batch_size,
           shuffle=False, num_workers=workers)
cifar_loader_train = DataLoader(cifar_train, batch_size=batch_size,
                               shuffle=True, num_workers=workers) cifar_loader_val =
DataLoader(cifar_val, batch_size=batch_size,
           shuffle=False, num_workers=workers) cifar_loader_test =
DataLoader(cifar_test, batch_size=batch_size,
           shuffle=False, num_workers=workers)

```

```

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This
DataLoader will create 4 worker processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to create. Please be aware that
excessive worker creation might get DataLoader running slow or even freeze, lower the worker numbe
er to avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(
import os
train_model = True
epochs = 50
epochs_nin = 100

lr = 0.004
lr_nin = 0.01
lr_scale = 0.5
momentum = 0.9

print_step = 5

deep_batch_size = 64
deep_num_classes = 10
deep_overshoot = 0.02
deep_max_iters = 50
    deep_args = [deep_batch_size, deep_num_classes, deep_overshoot,
                  deep_max_iters]

    if not os.path.isdir('weights/deepfool'):
        os.makedirs('weights/deepfool',
                    exist_ok=True)

    if not os.path.isdir('weights/fgsm'):
        os.makedirs('weights/fgsm', exist_ok=True)
# Загрузим и оценим стойкость модели Network-In-Network Model к FGSM и
DeepFool атакам на основе датасета CIFAR-10
fgsm_eps = 0.6
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth', map_location=
torch.device('cpu')))
evaluate_attack('mnist_lenet_fgsm.csv',
                'results', device, model, mnist_loader_test, mnist_min,
                mnist_max, fgsm_eps, is_fgsm=True)

print('')

evaluate_attack('mnist_lenet_deepfool.csv', 'results', device, model, mnist_loader_test,
mnist_min, mnist_max, deep_args, is_fgsm=False) if device.type == 'cuda':
    torch.cuda.empty_cache()
FGSM Test Error : 87.89% FGSM Robustness : 4.58e-01
FGSM Time (All Images) : 0.29 s FGSM Time (Per Image) : 28.86 us
DeepFool Test Error : 98.74% DeepFool Robustness : 9.64e-02
DeepFool Time (All Images) : 193.32 s DeepFool Time (Per Image) : 19.33 ms

# Загрузим и оценим стойкость модели LeNet к FGSM и DeepFool атакам на основе датасета
CIFAR-10
fgsm_eps = 0.2
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth',
map_location=torch.device('cpu')))
    evaluate_attack('mnist_fc_fgsm.csv', 'results', device, model,
                    mnist_loader_test, mnist_min, mnist_max, fgsm_eps,
is_fgsm=True) print('')

evaluate_attack('mnist_fc_deepfool.csv', 'results', device, model, mnist_loader_test,
mnist_min, mnist_max, deep_args, is_fgsm=False) if device.type == 'cuda':
    torch.cuda.empty_cache()
FGSM Test Error : 87.08% FGSM Robustness : 1.56e-01
FGSM Time (All Images) : 0.15 s FGSM Time (Per Image) : 14.99 us
DeepFool Test Error : 97.92% DeepFool Robustness : 6.78e-02
DeepFool Time (All Images) : 141.81 s DeepFool Time (Per Image) : 14.18 ms

# Загрузим и оценим стойкость модели LeNet к FGSM и DeepFool атакам на основе датасета
CIFAR-10
fgsm_eps = 0.1

```

```

model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth',
map_location=torch.device('cpu'))))
evaluate_attack('cifar_lenet_fgsm.csv', 'results', device, model, cifar_loader_test,
cifar_min, cifar_max, fgsm_eps, is_fgsm=True) print('')
evaluate_attack('cifar_lenet_deepfool.csv', 'results', device, model, cifar_loader_test,
cifar_min, cifar_max, deep_args, is_fgsm=False)

if device.type == 'cuda': torch.cuda.empty_cache()
FGSM Test Error : 91.71% FGSM Robustness : 8.90e-02
FGSM Time (All Images) : 0.40 s FGSM Time (Per Image) : 40.08 us
DeepFool Test Error : 87.81% DeepFool Robustness : 1.78e-02 DeepFool Time (All Images) : 73.27 s
DeepFool Time (Per Image) : 7.33 ms

```

## Выполним оценку атакующих примеров для сетей

*# LeNet на датасете MNIST*

```

fgsm_eps = 0.6
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth', map_location=device))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps,
deep_args,
                has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2, fig_width=25,
fig_height=11)

```

```

if device.type == 'cuda': torch.cuda.empty_cache()
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This
DataLoader will create 4 worker processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to create. Please be aware that
excessive worker creation might get DataLoader running slow or even freeze, lower the worker numbe
r to avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(

```

*# FCNet на датасете MNIST*

```

fgsm_eps = 0.2
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps,
deep_args, has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2, fig_width=25,
fig_height=11)

```

```

if device.type == 'cuda': torch.cuda.empty_cache()

```

In [19]:

*# Network-in-Network на датасете CIFAR*

```

fgsm_eps = 0.2
model = Net().to(device) model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))

```

```

display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps,
deep_args, has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2, fig_width=25,
fig_height=11, label_map=cifar_classes)

```

```

if device.type == 'cuda': torch.cuda.empty_cache()
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This
DataLoader will create 4 worker processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to create. Please be aware that
excessive worker creation might get DataLoader running slow or even freeze, lower the worker numbe
r to avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(

```

*# LeNet на датасете CIFAR*

```

fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))

```

```

display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps,
deep_args, has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2, fig_width=25,
fig_height=11, label_map=cifar_classes)

```

```

if device.type == 'cuda': torch.cuda.empty_cache()
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This
DataLoader will create 4 worker processes in total. Our suggested max number of worker in current

```

```
system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```

```
# Отразим отличия для fgsm_eps=(0.001, 0.02, 0.5, 0.9, 10) и выявим закономерность/обнаружим отсутствие влияние параметра eps для сетей FC LeNet на датасете MNIST, NiN LeNet на датасете CIFAR
```

```
fgsm_epss = [0.001, 0.02, 0.5, 0.9, 10]
```

```
for fgsm_eps in fgsm_epss:
```

```
    print(f"Используется fgsm_eps {fgsm_eps}") model = FC_500_150().to(device)
    model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
    display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max,
fgsm_eps, deep_args, has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2,
fig_width=25, fig_height=11)
```

```
    if device.type == 'cuda': torch.cuda.empty_cache() for fgsm_eps in fgsm_epss:
    print(f"Используется fgsm_eps {fgsm_eps}") model = FC_500_150().to(device)
    model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
evaluate_attack(f'mnist_fc_fgsm_eps{fgsm_eps}.csv', 'results', device, model,
mnist_loader_test, mnist_min, mnist_max, fgsm_eps, is_fgsm=True)
```

```
    if device.type == 'cuda': torch.cuda.empty_cache() for fgsm_eps in fgsm_epss:
    print(f"Используется fgsm_eps {fgsm_eps}") model = Net().to(device)
    model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))

    display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max,
fgsm_eps, deep_args, has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2,
fig_width=25, fig_height=11, label_map=cifar_classes)
```

```
    if device.type == 'cuda': torch.cuda.empty_cache() for fgsm_eps in fgsm_epss:
    print(f"Используется fgsm_eps {fgsm_eps}") model = Net().to(device)
    model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
    evaluate_attack(f'cifar_nin_fgsm_eps{fgsm_eps}.csv', 'results', device, model,
cifar_loader_test, cifar_min, cifar_max, fgsm_eps, is_fgsm=True)
```

```
    if device.type == 'cuda': torch.cuda.empty_cache()
```

```
Используется fgsm_eps 0.001
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: U
```

```
serWarning: This DataLoader will create 4 worker processes in total. Our sugg
```

```
ested max number of worker in current system is 2, which is smaller than what
```

```
this DataLoader is going to create. Please be aware that excessive worker cre
```

```
ation might get DataLoader running slow or even freeze, lower the worker numb
```

```
er to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```

```
Используется
fgsm_eps
0.02
```

```
Используется fgsm_eps 0.5
Используется fgsm_eps 0.9
Используется fgsm_eps 10
Используется fgsm_eps 0.001 FGSM Test Error : 3.07% FGSM Robustness : 8.08e-04
FGSM Time (All Images) : 0.72 s FGSM Time (Per Image) : 72.00 us Используется fgsm_eps 0.02
FGSM Test Error : 5.54% FGSM Robustness : 1.60e-02
FGSM Time (All Images) : 0.54 s FGSM Time (Per Image) : 53.68 us Используется fgsm_eps 0.5
FGSM Test Error : 99.21% FGSM Robustness : 3.86e-01
FGSM Time (All Images) : 0.56 s FGSM Time (Per Image) : 56.40 us Используется fgsm_eps 0.9
FGSM Test Error : 99.87% FGSM Robustness : 6.86e-01
FGSM Time (All Images) : 0.60 s FGSM Time (Per Image) : 60.29 us Используется fgsm_eps 10
FGSM Test Error : 99.87% FGSM Robustness : 1.47e+00
FGSM Time (All Images) : 0.55 s FGSM Time (Per Image) : 55.18 us Используется fgsm_eps 0.001
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This
DataLoader will create 4 worker processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to create. Please be aware that
excessive worker creation might get DataLoader running slow or even freeze, lower the worker numbe
r to avoid potential slowness/freeze if necessary.
warnings.warn(_create_warning_msg(

Используется fgsm_eps 0.02
Используется fgsm_eps 0.5
Используется fgsm_eps 0.9

Используется fgsm_eps 10
Используется fgsm_eps 0.001 FGSM Test Error : 10.12% FGSM Robustness : 8.92e-04
FGSM Time (All Images) : 1.09 s FGSM Time (Per Image) : 108.82 us Используется fgsm_eps 0.02
FGSM Test Error : 30.76% FGSM Robustness : 1.78e-02
FGSM Time (All Images) : 1.31 s FGSM Time (Per Image) : 131.25 us Используется fgsm_eps 0.5
FGSM Test Error : 82.65% FGSM Robustness : 4.40e-01
FGSM Time (All Images) : 1.12 s FGSM Time (Per Image) : 112.16 us Используется fgsm_eps 0.9
FGSM Test Error : 84.60% FGSM Robustness : 7.79e-01
FGSM Time (All Images) : 1.16 s FGSM Time (Per Image) : 116.42 us Используется fgsm_eps 10
FGSM Test Error : 87.53%
FGSM Robustness : 2.46e+00
FGSM Time (All Images) : 1.15 s FGSM Time (Per Image) : 115.36 us
import matplotlib.pyplot as plt fgsm_eps = [0.001, 0.02, 0.5, 0.9, 10]
fgsm_test_error_MNIST = [3.07, 5.54, 99.21, 99.87, 99.87]
fgsm_robustness_MNIST = [8.08e-04, 1.60e-02, 3.86e-01, 6.86e-01, 1.47e+00]
fgsm_test_error_CIFAR = [10.12, 30.76, 82.67, 84.62, 87.50] fgsm_robustness_CIFAR =
[8.92e-04, 1.78e-02, 4.40e-01, 7.79e-01, 2.46e+00]
plt.plot(fgsm_eps, fgsm_test_error_MNIST, label='MNIST') plt.plot(fgsm_eps,
fgsm_test_error_CIFAR, label='CIFAR')

plt.xlabel('FGSM Epsilon') plt.ylabel('Процент ошибок')
plt.title('Сравнение ошибки тестирования моделей') plt.legend()
plt.show()
```

In [24]:

```
plt.plot(fgsm_eps, fgsm_robustness_MNIST, label='MNIST') plt.plot(fgsm_eps,
fgsm_robustness_CIFAR, label='CIFAR')

plt.xlabel('FGSM Epsilon') plt.ylabel('Устойчивость') plt.title('Сравнение устойчивости
моделей') plt.legend()
plt.show()
```

По изображениям видно, что шум зависит от параметра `fgsm_eps`. Параметр `fgsm_eps` влияет на модель и ее защиту.

При увеличении `fgsm_eps` ошибка модели растет, что влияет на производительность.

Устойчивость модели увеличивается с `fgsm_eps`, но это делает модель более уязвимой к атакам.

Модель становится менее точной при обработке зашумленных данных.

Время выполнения не меняется с `fgsm_eps`, что показывает, что шум не влияет на скорость обработки.